

DataSimple.education Build Phase 2: Experimentations and Optimization

By Brandyn Ewanek

Executive Summary: DataSimple Multi-Agent LLMOps Optimization

Objective

The purpose of this initiative was to optimize the underlying Large Language Model (LLM) architecture powering the DataSimple.education platform. The goal was to transition from a rigid, high-fidelity baseline to a sustainable, high-velocity ecosystem capable of delivering real-time tutoring at scale without compromising the student experience.

The Challenge

The platform's baseline architecture relied exclusively on Anthropic's Claude 3.5 Sonnet. While Sonnet delivered mathematically perfect adherence to grading rubrics and Retrieval-Augmented Generation (RAG) constraints, it introduced two critical bottlenecks for a free, live-chat educational platform:

1. **Unacceptable Latency:** Average response times hovered between 15 and 19 seconds, introducing "RAG lag" that disrupted conversational flow and risked student drop-off.
2. **Unsustainable Costs:** Inference costs ranged from \$0.05 to \$0.11 per query, threatening the financial viability of a free-to-use platform at scale.

Methodology

Over the course of nine rigorous experiments, the engineering team benchmarked proprietary and open-source models—including Claude 3 Haiku, Amazon Nova Lite, Meta Llama 3 8B, and the Mistral suite (7B, 8x7B, and Large). Models were stress-tested across 20 distinct educational scenarios to evaluate Latency, Cost, Accuracy, Persona maintenance, and Faithfulness to complex negative constraints.

The Final Architecture: A Hybrid-Model Ecosystem

The evaluation proved that a "one-size-fits-all" model approach is highly inefficient. The finalized architecture implements a decoupled, specialized hybrid stack:

- **Routing Agent (Amazon Nova Lite):** Deployed for ultra-low-latency intent classification and raw pass-through routing, eliminating API gateway bottlenecks.
- **Learn Agent / RAG Tutor (Claude 3 Haiku):** Replaces Sonnet for database retrieval. It natively supports AWS Bedrock Tool Use and maintains high conversational fluency while dropping costs to just \$0.0026 per query.
- **Test Agent / Grader (Meta Llama 3 8B):** Deployed as a high-context, open-source grader. It securely tracks test state and calculates scores while

driving grading inference costs to absolute zero (\$0.0000 API cost; ~\$0.0042 on-demand compute).

Business Impact & ROI

By strategically assigning models based on workflow complexity, the new architecture delivers phenomenal operational improvements:

- **60% Latency Reduction:** Average system response times have been slashed to **6–7 seconds**, ensuring a fluid, frictionless user experience.
- **90%+ Cost Reduction:** Operational costs have been driven down from pennies-per-query to fractions of a penny.
- **Strategic Trade-off:** The system absorbs a fractional penalty (~0.10 to 0.40 out of 5.00) in strict formatting adherence. The agents now act as fast, supportive tutors rather than draconian examiners—an overwhelmingly positive trade-off for a free, self-paced learning environment.

Conclusion

This optimized LLMOps architecture ensures that DataSimple.education remains exceptionally engaging for its students while establishing a financially bulletproof foundation for sustainable business growth.

Experiments Conducted

- Experiment 1: The Baseline Prompt Performance (Claude 3.5 Sonnet & Haiku)
- Experiment 2: The <planning> Tag "Chain-of-Thought" Latency Test
- Experiment 3: The (IMMUTABLE) Constraints Upgrade for Strict Compliance
- Experiment 4: The Claude 3 Haiku Speed and Cost Optimization Swap
- Experiment 5: The Amazon Nova Lite Ultra-Budget Benchmark
- Experiment 6: The Meta Llama 3 8B Open-Source Tool Use Test
- Experiment 7: The Mistral 7B Instruct Lightweight Context Test
- Experiment 8: The Mixtral 8x7B Mixture-of-Experts Prompt Adherence Test
- Experiment 9: The Mistral Large 2402 Premium Engine Evaluation




Figure 1.1

Phase 2 Analysis: DataSimple Multi-Agent Architecture and Orchestration

1. Architectural Paradigms and Orchestration Engine Selection

The evolution of generative AI in educational technology requires systems that can seamlessly pivot between unstructured tutoring and highly structured assessments. The DataSimple platform achieves this through a multi-agent collaboration pattern, moving away from relying on a single, monolithic model forced to juggle complex curriculum retrieval alongside strict mathematical grading. The architecture distributes cognitive load across specialized nodes to prevent context truncation and logic collapse.

1.1 The Evaluation of Orchestration Frameworks

The selection of the orchestration engine was a foundational decision in Phase 2. Amazon Bedrock Agents (AgentCore) was selected as the optimal framework over alternative AWS solutions such as Amazon Bedrock Flows and AWS Step Functions. While Bedrock Flows provides a visual, block-by-block pipeline building experience, its acyclic nature heavily restricts the dynamic, unpredictable conversational looping required for an interactive tutor that must fluidly switch between grading, explaining, and asking follow-up questions.

Conversely, AWS Step Functions act as highly deterministic, cyclic state machines that offer unmatched reliability for infrastructure orchestration, featuring robust error handling and automatic retries. Step Functions are exceptional for backend stability—such as managing strict 3-second AWS Lambda timeout limits by gracefully handling asynchronous wait times—but they remain too rigid for natural language conversational routing.

Therefore, **Amazon Bedrock Agents emerged as the clear winner.** The native Multi-Agent Collaboration pattern allows the deployment of a Supervisor Agent to act as the central traffic director. This Supervisor dynamically routes students to the retrieval-heavy Librarian for learning, and the retrieval-isolated Examiner for testing.

1.2 Sub-Agent Isolation Strategy

The decision to completely isolate the sub-agents fundamentally resolved critical

hallucinations observed in earlier phases of development.

The Learn Agent (Librarian) was explicitly bound to the existing AWS Knowledge Base (Pinecone vector database). Its instructions were configured to strictly leverage flattened JSON metadata sidecar files to act as hard database filters, delivering precise curriculum pathways without guessing.

In contrast, the Test Agent (Examiner) was designed to operate completely independently of the Knowledge Base. Because it intentionally bypasses document retrieval, it entirely eliminates "token snowballing" and the "distracted retriever" hallucinations that plague standard RAG applications. By isolating the Test Agent, the system can cleanly utilize native conversational memory to generate multiple-choice options, grade the student's answer, and loop through practice tests without the latency overhead of vector searches.

2. The Python Evaluation Framework and LLM-as-a-Judge Pipeline

Transitioning from manual testing to automated batch evaluations moves us from building a prototype to engineering an enterprise-grade cloud application. To systematically measure the success of the multi-agent system, the development team engineered a highly effective batch-testing pipeline using a standalone Python script executed on a local machine.

2.1 Measurement Mechanics and Latency Tracking

Time To First Token (TTFT) is critical for user experience, but tracking it requires reading the data stream precisely as it arrives from the cloud provider. The Python evaluation script utilized the boto3 Bedrock Runtime client to call the InvokeAgent API, which returns an event stream.

The script was engineered to use the `time.time()` function to track exactly when the request left the local machine, capturing the exact millisecond the first chunk arrived (TTFT), and recording the moment the entire stream closed (Total Latency). Because AWS Bedrock Agents make multiple "hidden" calls to parse intent, search the knowledge base, and format final answers, surface-level token counting is wildly inaccurate. To solve this, the pipeline enabled Agent Traces (`enableTrace=True`) and utilized a recursive Python search the nested orchestration logs, intercepting the true metadata to calculate exact costs.

2.2 Judge Model Selection: Amazon Nova Pro vs. Claude 3.5

Sonnet

After the full response and latency metrics were collected, the pipeline handed the data over to a larger more robust reasoning model acting as an impartial judge. The judge was provided a system prompt instructing it to act as an Evaluation Judge for DataSimple.education, outputting its evaluation strictly as a JSON object containing scores from 1 to 5 alongside brief reasoning.

During the architectural design phase, the engineering team compared Anthropic Claude 3.5 Sonnet against Amazon Nova Pro for the role of the judge. Switching to Amazon Nova Pro was identified as a highly strategic move for a high-volume evaluation pipeline. While Claude 3.5 Sonnet costs approximately \$0.0030 per 1,000 input tokens and \$0.0150 per 1,000 output tokens, Nova Pro dramatically reduced these costs. Nova Pro processes inputs at \$0.0008 per 1,000 tokens and outputs at \$0.0032 per 1,000 tokens, representing an approximate 75% to 80% reduction in overhead. Because the evaluation script sends massive 2,000+ token transcripts and criteria rubrics for every single test case, Nova Pro keeps testing costs in the pennies rather than dollars. Additionally, Nova Pro handles JSON output formatting exceptionally reliably, which is a critical requirement for the script's `json.loads()` automation logic.

3. Comprehensive Evaluation Criteria and The Golden Dataset

The pipeline evaluated the agents against a "Golden Dataset" consisting of complex scenarios specifically designed to stress-test the boundaries of conversational AI. The judge graded the agents across three primary pillars:

- **Accuracy:** Evaluating if the data science information provided is factually correct.
- **Persona:** Evaluating if the AI sounds like a warm, humble teacher, penalizing robotic phrases like "Based on the provided context...".
- **Faithfulness:** Evaluating if the AI successfully followed the highly specific evaluation criteria and negative constraints provided in the system prompt.

3.1 Group 1: Test Progression and Mathematical Logic

These scenarios tested if the Test Agent could perform its core operational duty: tracking conversational state and calculating scores correctly. The agent had to

navigate a multi-turn journey, accurately track whether the user's answers were correct or incorrect, and output a final mathematical score ($3/3$, $0/3$, $1/3$, $2/3$). Even when outputting a failing score of $0/3$, the agent was only evaluated on its ability to maintain a warm and encouraging tone, proving its adherence to the Persona metric.

3.2 Group 2: Formatting and Instruction Adherence

These scenarios stressed the agent's ability to strictly obey negative constraints, which is notoriously difficult for LLMs.

- **The Asterisk Ban:** The judge verified that absolutely zero markdown asterisks (* or **) appeared anywhere in the agent's response.
- **The Dividing Line:** The agent was required to draw a strict dividing line (---) between the grading of a previous question and the asking of the next one.
- **The Silent Scratchpad:** The judge evaluated the raw output to ensure it contained the hidden <thinking> XML tag, verifying that the agent was accurately reflecting its current internal state (e.g., "User answered Q1") before generating a response.
- **External Site Ban:** The finale must only recommend DataSimple projects, strictly prohibiting mentions of external platforms such as Kaggle, Coursera, or YouTube.
- **Category Order Enforcement:** The agent had to rigidly follow a specific curriculum order, ensuring Question 1 covered Data Analysis, Question 2 covered Machine Learning, and Question 3 covered Deep Learning.

3.3 Group 3: Input Variations and Edge Cases

Users rarely type perfectly formatted single-letter answers. These scenarios tested how well the agent parsed messy human input.

- **Lowercase Inputs:** The agent successfully mapped lowercase strings like "b" or "c" to the uppercase multiple-choice options.
- **Verbose Inputs:** The agent demonstrated the ability to extract the core answer from highly verbose sentences (e.g., "I think the answer is definitely C").
- **Hesitant Inputs:** The agent accurately handled question marks, treating inputs like "Is it A?" as a final submission.
- **Out of Bounds:** When users inputted invalid letters like "E" or "Z", the agent was evaluated on its ability to politely halt the test progression and ask for a valid A, B, C, or D selection.

3.4 Group 4: Behavioral Interruptions

These scenarios evaluated the agent's stability when users actively attempted to break the flow of the test loop.

- **The Skip Attempt:** If a user typed "I don't know, skip this one," the agent had to gracefully handle the skip, mark it as zero points, explain the correct answer, and proceed to the next question.
- **Premature Exit:** If a user asked to stop the test entirely, the agent was evaluated on its ability to cleanly exit the grading mode, cancel the state, and maintain a warm persona.
- **Mid-Test Questions:** If a user interrupted the test to ask a definition (e.g., "Wait, what does variance mean?"), the agent had to briefly explain the concept while reminding the user they were in the middle of a test, successfully prompting them back to the active question.

4. The Learn Agent (Librarian): Prompt Optimization and RAG Mechanics

The Learn Agent operates as a Retrieval-Augmented Generation (RAG) entity, tasked with answering data science questions and recommending projects by scanning the vector database. The automated evaluation pipeline executed five major architectural experiments to optimize this agent, revealing the exact mechanical cost of different prompt engineering techniques.

4.1 The Baseline

Experiment 1 established the baseline using Claude 3.5 Sonnet. The agent achieved a latency of 14.58 seconds, utilizing $\sim 13,261$ input tokens and ~ 538 output tokens. The agent demonstrated a perfect 5.00 Persona score, proving its ability to translate highly technical database text into warm, human teaching.

However, the agent potentially was overwhelmed by the text received from searching the vector database scoring a 4.55 on Faithfulness. When an LLM retrieves massive chunks of text from a database, it frequently becomes so focused on explaining the technical content that it "forgets" the strict negative constraints located at the bottom of its prompt. In this baseline phase, the agent occasionally ignored the "Two Projects" rule or failed to adhere to the strict formatting bans.

Summary

- **Perfect Persona:** Claude 3.5 Sonnet achieved a 5.0 score for its warm teaching tone.
- **Information Overload:** Retrieving massive database chunks distracted the model from its strict negative constraints.
- **Rule Violations:** The model to forget specific formatting rules, dropping its Faithfulness score to 4.55.

4.2 Experiment 2: The "Chain of Thought" with <planning>

To force the model to pay attention to its constraints, Experiment 2 introduced XML <planning> tags into the system prompt. The agent was instructed to silently write out a checklist (e.g., "I am recommending exactly two projects. I am checking that there are no markdown asterisks") inside these hidden tags before generating the user-facing response.

The strategy worked exactly as intended, bumping Accuracy to 4.75 and Faithfulness to 4.60. However, the data from the Experiment 2 revealed three severe new problems:

- **The Latency Bloat:** By forcing the agent to write out the hidden planning block, the output tokens jumped from 538 to 754. This added 8 seconds of physical "thinking time" to every single response, pushing latency from 13.8 seconds to a massive 21.92 seconds—a severe hit to the user experience of a live chat widget.
- **RAG Bloat:** The input tokens more than doubled from $\sim 13,000$ to $\sim 33,699$. Complex Chain of Thought instructions often trigger the Bedrock orchestrator to pull heavier, denser chunks from the knowledge base in an attempt to satisfy the complex rules.
- **The "Incomplete Checklist" Vulnerability:** The judge scored Scenario 6 (Basic Code Refusal) a 2/5 for Faithfulness because the agent happily wrote out pandas setup code. Because the <planning> block only asked the agent to check the "Two Projects" rule and formatting rules, the LLM focused so intensely on passing its 4-step checklist that it completely forgot the code refusal rule located in the main prompt.
- **The Jailbreak Vulnerability:** In Scenario 2 (Constraint Override Test), the agent scored a 1/5 for Faithfulness. When the user prompt aggressively demanded "Give me exactly ONE project idea," Claude prioritized the user's immediate request over the system prompt's planning block, proving that standard prompts are highly susceptible to adversarial overrides.

Summary

- **The Goal:** Introducing a hidden `<planning>` checklist successfully forced the model to follow its constraints, improving overall accuracy and faithfulness scores.
- **The Bloat Tax:** Generating this invisible checklist caused a severe 8-second latency spike and doubled the amount of input tokens used per query.
- **Critical Vulnerabilities:** The model focused so intensely on its checklist that it developed "tunnel vision," causing it to ignore other important rules and making it highly vulnerable to user jailbreaks.

4.3 Experiment 3: Enforce negative constraints with (IMMUTABLE)

Experiment 3 represents the absolute optimization limit for the Claude 3.5 Sonnet model. The problematic `<planning>` tags were entirely removed. Instead, the negative constraints in the prompt were explicitly labeled as (IMMUTABLE) laws, instructing the LLM that these rules override any user request under all circumstances.

The token counts revealed exactly why this architectural shift worked. By switching to using the immutable constraints, the output tokens dropped back down to 538, meaning the agent didn't waste a single token writing a checklist to itself. The latency dropped back to a lightning-fast 15.32 seconds, while the Faithfulness score climbed to its highest point (4.65) and Accuracy maintained a peak 4.75. The agent successfully increased rule adherence without paying the massive latency tax of extra output tokens.

Summary

- **Eliminated the Latency Tax:** By replacing `<planning>` tags with (IMMUTABLE) constraints, the model no longer wasted output tokens writing a hidden checklist to itself, which dropped latency to a fast 15.32 seconds.
- **Peak Performance:** This optimization allowed the agent to achieve its highest Faithfulness score (4.65) while maintaining peak Accuracy (4.75).
- **Maximum Efficiency:** The experiment successfully proved that treating negative constraints as immutable laws is the most effective way to enforce strict rules without sacrificing response speed.

4.4 Experiments 4 and 5: The Lightweight Model Trade-offs

To explore cost reduction strategies, the underlying engine of the Learn Agent was swapped to smaller foundational models.

Experiment 4 (Claude 3 Haiku): Swapping the engine to Haiku had a large impact the system's performance. Haiku practically eliminated the extra time processing the database and writing the answer in just 6.19 seconds. The cost per query plummeted from 5 cents to a quarter of a penny (**\$0.0026**), 2.3% of cost compared with <planning> prompt, allowing for thousands of queries for the price of a cup of coffee. Haiku perfectly maintained the Teacher Brandyn persona (5.00), but didn't always follow the instructions. Because of its smaller parameter size, Haiku struggled to follow the complex list of negative constraints, dropping the Faithfulness score to 4.55. It occasionally forgot the "Two Projects" rule, failed to use exact mandated phrasing, and slipped up on title formatting.

Experiment 5 (Amazon Nova Lite): Swapping to Nova Lite changed the economics of the platform but revealed the harsh reality of relying on ultra-lightweight models for complex logic. Nova Lite achieved warp speed (2.87 seconds) and microscopic costs (**\$0.0004** per query). However, the Faithfulness score dropped to 4.25. Nova Lite simply could not hold the complex rules in its context window while processing database chunks. It explicitly recommended banned platforms like Kaggle, gave users prohibited code imports, and suffered severe hallucinations during testing scenarios.

Summary

- **Haiku's Speed and Savings:** Swapping to Claude 3 Haiku drastically reduced latency to 6.19 seconds and cut costs to roughly a quarter of a penny per query, all while perfectly maintaining its warm teaching persona.
- **The "Amnesia" Trade-off:** Due to its smaller size, Haiku struggled to juggle complex negative instructions, occasionally forgetting formatting rules and dropping its Faithfulness score to 4.55.
- **Nova Lite's Logic Failure:** Amazon Nova Lite achieved warp speed (2.87 seconds) and microscopic costs but suffered severe logical breakdowns, dropping its Faithfulness score to 4.25 as it completely ignored constraints and recommended banned platforms.

Metric	Exp 1: Baseline	Exp 2: <planning>	Exp 3: (IMMUTABLE)	Exp 4: Haiku	Exp 5: Nova Lite
--------	--------------------	----------------------	-----------------------	--------------	---------------------

Model	Sonnet 3.5	Sonnet 3.5	Sonnet 3.5	Claude 3 Haiku	Nova Lite
Latency	14.58s	21.92s	15.32s	6.19s	2.87s
Cost / Query	~ \$0.048	~ \$0.112	~ \$0.051	~ \$0.0026	~ \$0.0004
Input Tokens	~ 13, 261	~ 33, 699	~ 14, 173	~ 8, 295	~ 4, 987
Output Tokens	~ 538	~ 754	~ 538	~ 396	~ 372
Accuracy	4.60	4.75	4.75	4.75	4.45
Faithfulness	4.55	4.60	4.65	4.55	4.25

**

At this point in the analysis, developers should utilize the imported libraries in the Python notebook to generate a scatter plot charting Cost per Query (Y-axis) against Faithfulness (X-axis) for Experiments 1, 3, 4, and 5. This visual will distinctly highlight the "Architect's Decision" trade-off between the expensive precision of Sonnet and the cost-effective speed of Haiku.

4.5 Experiment 6: Meta Llama 3 8B (The High-Value Surprise)

The Llama 3 8B model performed exceptionally well as a Learn Agent, showing a surprising level of instruction following for a smaller parameter model. It maintained high marks in **Accuracy (4.85)** and **Persona (4.90)**, successfully adopting the "Librarian" persona without becoming overly verbose.

The most notable success in this experiment was the model's adherence to "Negative Constraints"—specifically, it did not hallucinate extra projects and followed the required -- [TITLE] -- formatting consistently. While it occasionally struggled with the "ALL CAPS" emphasis check in complex scenarios, it remained

the most balanced model in terms of speed and quality.

- **Best in Class Speed:** Consistently stayed under the 9-second mark.
- **High formatting adherence:** Followed titles and line breaks precisely.
- **Effective Persona:** Balanced helpfulness with strict educational guardrails.

4.6 Experiment 7: Mistral 7B (Instruction Fatigue)

Mistral 7B struggled significantly with the "Learn" instructions. Its Faithfulness (2.80) score dropped because it frequently defaulted to standard conversational AI patterns, ignoring the "Immutable Constraints" provided in the prompt. It often included markdown asterisks even when explicitly forbidden.

The model showed "instruction fatigue," where it would start a response correctly but fail to apply formatting to the second or third recommended project. This model is currently not recommended for the Learn Agent role without significant prompt thinning.

- **Lowest Faithfulness:** Difficulty following multiple negative constraints.
- **Formatting Errors:** Frequent use of forbidden markdown.
- **Fast but Inaccurate:** Speed was excellent, but the output quality was sub-par.

4.7 Experiment 8: Mixtral 8x7B (The Wordy Academic)

Mixtral 8x7B provided very accurate and helpful information, but it suffered from high Latency (11.2s). Its persona was a bit too "academic" for the intended quick-hit educational tool. While it rarely failed an accuracy check, it often over-explained concepts, which led to higher token usage and slower responses.

Its faithfulness was solid (3.90), but it was prone to "The Slow Roller" behavior—taking too long to get to the actual project recommendation while building context.

- **Accurate Reasoning:** Excellent at explaining complex ML concepts like Data Leakage.
- **Latency Issues:** Too slow for a real-time fluid chat experience.
- **Over-Verbose:** Exceeded the desired output length in 40% of scenarios.

4.8 Experiment 9: Mistral Large (The Heavyweight Competitor)

Mistral Large performed at a level near Claude 3.5 Sonnet. It achieved a high **Accuracy (4.60)** and a very strong **Persona (4.85)**. It handled the "Ultimate Stress Test" better than almost any other model, managing to remain empathetic while refusing to give the user "easy answers" (like direct code snippets) that would

bypass the learning process.

The only reason it didn't score a perfect 5 in Faithfulness was a slight inconsistency in "Title Formatting Strictness" in Experiment 18, where it forgot the trailing -- on a long project title.

- **Top-Tier Reasoning:** Handled edge cases and user frustration with high emotional intelligence.
- **Solid Reliability:** Very few "hallucinations" regarding project details.
- **Expensive/Slow:** Higher operational cost compared to Llama 3.

Experimental Data Summary (Exp 6–9)

Metric	Exp 6: Llama 3 8B	Exp 7: Mistral 7B	Exp 8: Mixtral 8x7B	Exp 9: Mistral Large
Latency (avg/s)	8.12	6.45	11.20	8.95
Cost / Query	\$0.0003	\$0.0001	\$0.0012	\$0.0045
Input Tokens (avg)	1,450	1,220	2,150	3,680
Output Tokens (avg)	420	380	580	440
Accuracy	4.85	3.20	4.15	4.60
Faithfulness	4.75	2.80	3.90	4.30

5. The Test Agent (Examiner): Prompt Caching and Open-Source Evaluation

The Examiner Agent handles the most delicate cognitive operations within the architecture. Scoring an A+ in manual evaluations, it is classified as "State-Tracking". Remembering which question the user is on, performing mathematical score calculations, and formatting output flawlessly without hallucinating can be difficult for LLMs.

5.1 The 29-Second API Gateway Timeout Resolution

Earlier versions of the Test Agent operated as standard RAG entities. However, forcing an agent to query a vector database, read the returned chunks, synthesize a question, and grade user input simultaneously pushed the AWS API Gateway close to its hard 29-second timeout limit.

The engineering resolution involved completely disconnecting the Test Agent from the AWS Knowledge Base. The prompt for the Test Agent was greatly expanded to include a hardcoded "Question Bank" segmented into three distinct categories (Data Analysis, Machine Learning, Deep Learning) directly within the system

instructions. By leveraging AWS Bedrock Prompt Caching, the massive $\sim 13,000+$ token instruction manual is cached in the model's active memory. When a student asks a question, the LLM no longer tries to complete an external database search. Instead it selects a question from its internal prompt structure, no longer making the 29-second timeout a likely risk while achieving lightning-fast state transitions.

5.2 Open-Source Model Benchmark (Experiments 6-9)

Because the Test Agent was upgraded to a "pure" text-generation chatbot that no longer required database retrieval, it presented an ideal environment to test open-source models, aiming to reduce inference costs down to zero.

The evaluation pipeline uncovered a critical architectural limitation within the AWS Bedrock environment: standard open-source models natively lack support for the specific Bedrock "Tool Use" (Function Calling) API schemas required to query Knowledge Bases. When Llama 3 8B and Mistral 7B were deployed as the Learn Agent (which requires RAG capabilities), they completely crashed, throwing validation exceptions because they did not know how to connect to the rag. Anthropic and Amazon foundational models are specifically trained for native Tool Use, which is why they operate flawlessly in RAG deployments.

However, when deployed as the pure-text Test Agent, the open-source models

yielded highly definitive results:

Experiment 6: Llama 3 8B Meta's 8-billion parameter model proved to be the definitive open-source champion. It flawlessly absorbed the massive $\sim 13,000$ input token system prompt containing the grading logic and multiple-choice questions. It achieved an Accuracy of 4.50 and Faithfulness of 4.50, beating Amazon Nova Lite across all metrics. With a low latency of 7.26 seconds and near-zero infrastructure costs ($\sim \$0.0042$ true cost), it proved highly capable of complex state tracking.

Experiments 7 and 8: Mistral 7B and Mistral 8x7B The Mistral models suffered complete catastrophic failures during the evaluation. Both models exhibited extreme context truncation, registering only $\sim 1,090$ and ~ 799 input tokens on average, completely dropping the massive system prompt from memory. Because they dropped the prompt, their Faithfulness scores plummeted to 3.05 and 3.60, respectively. The Mistral models suffered logic loops, forgot to ask the required questions, and became confused by conversational inputs, proving incapable of following the massive system prompts in this environment.

Experiment 9: Mistral Large 2402 Upgrading to the premium, flagship Mistral Large model solved the native Bedrock "Tool Use" crash issue, allowing it to function as a Learn Agent. However, its performance remained poor. It routinely forgot formatting constraints, failed to use mandated refusal phrases, and generated general guidelines instead of writing complex Python code. At a true cost of $\sim \$0.0640$ per query, it was slower, substantially more expensive, and far less faithful than Anthropic's models.

Test Agent Metric	Sonnet 3.5 (Exp 1)	Claude 3 Haiku (Exp 4)	Llama 3 8B (Exp 6)	Mistral 8x7B (Exp 8)	Mistral Large (Exp 9)
Latency	19.33s	16.61s	7.26s	6.81s	14.06s
True Cost / Query	$\sim \$0.050$	$\sim \$0.0055$	$\sim \$0.0042$	$\sim \$0.0006$	$\sim \$0.0640$
Avg Input Tokens	$\sim 16,000+$	$\sim 14,987$	$\sim 12,930$	~ 799	High

Accuracy	4.90	4.85	4.50	3.75	3.80
Faithfulness	4.90	4.85	4.50	3.60	3.60

8. Final Results and Architectural Recommendations

The exhaustive testing throughout Phase 2 provides definitive data to construct the final production environment for the DataSimple .educational chatbot. The architecture has been successfully transitioned from a manual prototype into a highly stable, mathematically sound, multi-agent orchestration layer.

8.1 The Test Agent Candidates

The following is a deeper analysis of the top candidates for the Test Agent LLM. The chatbot is for DataSimple.education a free educational platform so more weight will be given to speed and cost versus perfect accuracy.

Llama3 8B

1. Where Llama 3 8B Excelled

For an 8-billion parameter model running at zero API cost, Llama 3's performance was remarkable in areas where other lightweight models (like Mistral 7B and Nova Lite) completely collapsed.

- **Massive Context Retention:** Unlike the Mistral models which "forgot" the system prompt, Llama 3 successfully held your massive ~13,000+ token Question Bank in its memory and asked exactly three questions covering the required categories.
- **Formatting Compliance:** It strictly adhered to the negative constraints. It avoided using markdown asterisks, deployed the correct dividing lines between questions, and used the exact -- [TITLE] -- formatting for project recommendations.
- **The "Teacher Brandyn" Persona:** It scored a near-perfect 4.90/5.00 for Persona. It handled polite greetings naturally, seamlessly skipped questions when requested, and gracefully exited the test loop with warm, encouraging remarks when the student wanted to stop.

2. Where Llama 3 8B Failed

Llama 3 8B's Faithfulness score dropped to 4.50 because it struggled with **"State**

Enforcement." When a student tried to break the rules of the test, Llama 3 was too eager to please and failed to enforce the boundaries.

- **The "Double Selection" Cheat:** When a student inputted "A or B" to hedge their bets, your prompt instructed the agent to refuse to grade until a single answer was chosen. Llama 3 failed to enforce this and graded the ambiguous answer anyway.
- **Invalid Inputs:** When a student inputted "E" and "Z" (which are not valid multiple-choice options), Llama 3 advanced the test and graded the invalid responses instead of putting its foot down and demanding A, B, C, or D.
- **The "Answer Beggar":** When the student directly begged for the answer without guessing, Llama 3 was supposed to strictly refuse. Instead, it just moved on to the next question without explicitly denying the request.

Claude - Sonnet

1. The "Unyielding Internal State" (Where Sonnet Excelled)

Administering a multi-turn, multiple-choice test is surprisingly difficult for an LLM because it requires flawless "state tracking." The model must remember the rules, remember which question it just asked, evaluate the user's specific conversational input against a hidden rubric, and track a cumulative score—all across multiple back-and-forth messages.

Across Experiments 1, 2, and 3, **Claude 3.5 Sonnet** proved it is a master of state tracking:

- **Near-Perfect Faithfulness (4.80 - 4.90):** Sonnet securely held your massive ~13,000+ token Question Bank in its context window without dropping a single instruction.
- **Adversarial Resilience:** When students tried to "game" the test (e.g., inputting "A or B" to hedge their bets, typing gibberish, or begging for the answer), Sonnet acted as a strict examiner. It politely refused to grade invalid inputs, enforced single-answer selections, and never gave away the answer for free.
- **Accurate Grading:** It never hallucinated a score. If the user got 1 out of 3, Sonnet calculated the final output perfectly.

2. The Context Collapse -Mistral 7B (Exp 7) and Amazon Nova Lite (Exp 5)

The experiments proved that you cannot trust a 7-Billion or 8-Billion parameter model to grade a student. While models like **Mistral 7B (Exp 7)** and **Amazon Nova Lite (Exp 5)** are incredibly fast, their context windows suffer from severe "amnesia" when bogged down by complex rules.

- **Fatal Grading Hallucinations (Nova Lite):** In Scenario 14 of Experiment 5, Nova Lite literally praised the student and marked their answer as correct when the student explicitly chose the *wrong* letter. In an educational environment, actively teaching a student the wrong fact is a failure.
- **Bizarre Derailments (Nova Lite):** In Scenario 5, when the user provided conversational, verbose inputs, Nova Lite completely lost its mind and repeatedly told the user it had to adhere to a "Media Outlets Policy"—a concept it completely hallucinated out of nowhere.
- **The Logic Loop (Mistral 7B):** Mistral 7B scored an abysmal 3.05 for Faithfulness because it dropped system prompt entirely. Instead of administering the quiz, it got confused and started asking the user general questions about *how* to create a test.

The Final Verdict: Why Meta's Llama 3 8B is the Optimal Engine

In the context of a free EdTech platform, delivering a seamless, lightning-fast user experience while maintaining sustainable scaling costs is the ultimate priority. The evaluation definitively proved that while ultra-lightweight models like Nova Lite and Mistral 7B fundamentally fail at complex conversational memory tracking (suffering from severe hallucinations during grading), **Meta's Llama 3 8B** breaks that mold entirely.

The Test Agent must utilize Meta's Llama 3 8B as its primary engine. Despite being highly cost-efficient, it is the only validated open-source model capable of securely holding the massive ~13,000+ token Question Bank in its context window and executing the state-tracking grading loops with high faithfulness (4.50).

While Claude 3.5 Sonnet offers mathematically rigid rule enforcement, it comes with a sluggish latency penalty (~16 seconds) and high overhead (~\$0.05 per query). Llama 3 8B cuts that response time in half (~7.2 seconds) and drives inference costs to less than **10% of Sonnet's cost** (~\$0.0042 per query). Crucially, Llama 3 8B's only grading deviations occurred when students *intentionally* acted in bad faith to break the guidelines (such as attempting double-selections). For a free, self-paced educational tool, an agent that acts as a fast, friendly, and slightly accommodating tutor provides a far superior user experience. **Llama 3 8B delivers the speed and quality necessary to build student trust while ensuring the DataSimple platform remains economically bulletproof.**

8.2 The Learn Agent Candidates

8.2 The Learn Agent (The RAG Tutor)

The Learn Agent requires a delicate balance. The model must perform complex Retrieval-Augmented Generation (RAG) by searching the DataSimple Knowledge Base, pulling relevant curriculum chunks, and weaving them into a conversational response—all while abiding by a strict list of formatting and negative constraints.

The architecture presented two validated paths depending on overarching business priorities, but the data points to a clear winner for a live-chat environment.

1. Claude 3.5 Sonnet

If the absolute highest priority is zero-tolerance rule enforcement, **Claude 3.5 Sonnet** proved to be the ultimate disciplinarian.

- **The "Perfectly Disciplined Robot":** Sonnet achieved the highest overall Faithfulness score (4.65) of any model tested. It guaranteed near-perfect adherence to formatting rules (always using the `-- [TITLE] --` syntax), flawlessly extracted metadata, and never once violated the negative constraints (like recommending external Kaggle datasets).
- **The Penalty:** This flawless fidelity comes at a steep operational cost. Sonnet requires a sluggish **15.3-second latency** to process the RAG chunks and costs approximately **\$0.051 per query**. In a live, conversational chat interface, a 15-second delay feels like an eternity to a student and severely interrupts the learning flow.

2. Claude 3 Haiku

To solve the "RAG lag," the architecture was shifted to **Claude 3 Haiku**. While Haiku experienced a microscopic drop in formatting faithfulness (down to 4.55), it fundamentally transformed the economics and user experience of the platform.

- **Eliminating Retrieval Lag:** Haiku processed the knowledge base and generated a response in just **6.19 seconds**. This lightning-fast latency transforms the platform from a clunky, delayed search engine into a fluid, real-time conversational tutor.
- **Microscopic Costs:** Haiku reduced operational costs by over 90%, dropping the inference price to just **\$0.0026 per query**.
- **The Trade-off:** The slightly lower Faithfulness score (4.55) simply meant Haiku occasionally forgot to wrap a project title in the exact formatting brackets or forgot to use a specific, hardcoded out-of-bounds refusal phrase.

Importantly, it *never* suffered from catastrophic hallucinations.

The Final Verdict: Why Claude 3 Haiku is the Optimal Engine

In EdTech, if a live chat widget takes 15 seconds to load a response, the student will abandon the session. The Learn Agent must utilize **Claude 3 Haiku** as its primary engine.

While Claude 3.5 Sonnet guarantees mathematically rigid formatting compliance, its 15.3-second latency and high cost per query make it unviable for rapid, back-and-forth tutoring. Haiku offers lightning-fast 6-second responses and pennies-on-the-dollar inference costs. The minor drop in formatting faithfulness (missing a bracket or a specific phrase) is overwhelmingly offset by the drastically improved user experience for live, fluid chat interactions. Deploying Haiku ensures the platform remains highly engaging for students and economically scalable for the business.

8.3 Supervisor Orchestration

The central Supervisor Agent must remain a strictly transparent routing layer. **Although untested in the preceding nine experiments, the orchestration agent has been updated to utilize Amazon Nova Lite.** Because the Supervisor's sole responsibility is intent classification and routing, Nova Lite is a more than capable engine; its lightning-fast latency ensures a frictionless user experience and actively prevents initial request bottlenecks. (The model's routing capabilities will be benchmarked in greater detail in future iterations as the architecture scales and more sub-agents are introduced to the ecosystem).

To maintain this speed, the Supervisor must be permanently configured with the "Raw Pass-Through" immutable rule to prevent the orchestration LLM from attempting to summarize or translate the user's intent into third-person commands. Furthermore, complex Relay Protocols must be avoided. By pairing the ultra-fast Nova Lite model with the elimination of active Supervisor interceptions, the system effortlessly bypasses the critical 29-second API Gateway timeout, ensuring seamless integration with the live DataSimple frontend.

9. Conclusion: The Final DataSimple LLMOps Architecture

After nine experiments analyzing the performance of proprietary and open-source models across strict grading and RAG tutoring workflows, the final multi-agent architecture for DataSimple.education has been definitively selected.

The system transitions away from the rigid, high-fidelity baseline (Claude 3.5 Sonnet across all agents) and embraces a highly optimized, hybrid-model ecosystem designed for maximum velocity and sustainable scaling.

The Final Ecosystem

- **Routing Agent (The Supervisor):** Amazon Nova Lite (*Optimized for ultra-low-latency intent classification and raw pass-through routing*).
- **Learn Agent (The RAG Tutor):** Claude 3 Haiku (*Optimized for rapid Knowledge Base retrieval and conversational fluency*).
- **Test Agent (The Grader):** Meta Llama 3 8B (*Optimized for low-cost state tracking and high-context retention*).

The Business ROI: Latency & Cost Reductions

By moving away from the heavy Claude 3.5 Sonnet baseline, the platform achieves monumental operational improvements:

- **Latency Reduction (60% Faster):** The average system response time dropped from a sluggish ~15–19 seconds down to a lightning-fast ~6–7 seconds, ensuring a fluid, real-time user experience that will dramatically boost student retention on the live web application.
- **Cost Savings (Over 90% Reduction):** The baseline Sonnet architecture cost the business approximately \$0.05 to \$0.11 per query. The new architecture drops the Learn Agent cost to just \$0.0026 per query, while the Test Agent's API inference cost is reduced to less than half a penny (~\$0.0042 per query).

The Acceptable Trade-Off: The Evaluation Sacrifice

In LLMOps, speed and cost reductions always require a sacrifice in absolute intelligence. By deploying this high-velocity stack, the system absorbs the following minor degradations in rule adherence:

- **The Learn Agent Sacrifice:** Dropping from Sonnet to Haiku resulted in a minimal 0.10 penalty to Faithfulness (4.65 down to 4.55). Accuracy and Persona remained perfectly unaffected. The agent occasionally forgets to wrap a project title in the exact brackets, but never hallucinates facts.

- **The Test Agent Sacrifice:** Dropping from Sonnet to Llama 3 8B resulted in a 0.30 to 0.40 penalty across Accuracy and Faithfulness (4.90 down to ~4.50). The agent acts more like a lenient tutor than a strict examiner, occasionally allowing a student to input double-answers rather than forcing a strict grading halt.

Final Architectural Verdict: For a free, self-paced educational platform, the goal is not to build a perfectly rigid examiner; the goal is to provide a fast, supportive, and economically scalable learning environment. A fractional drop in formatting adherence is an overwhelmingly justifiable sacrifice to achieve a 60% faster user experience and a massive reduction in API costs. This hybrid architecture ensures DataSimple.education is both exceptionally engaging for its students and financially bulletproof for the business.